

Detecting Microaggressions Using LLM-Generated Counterfactuals: Transformer Cross-Domain Evaluation and LLM Classification

Nura Hossainzadeh, Neeharika Kotte, Carlos Schrupp

UC Berkeley



Introduction & Background

- “**Microaggression**”: coined by Chester Pierce in 1970
- “**Brief and commonplace indignities** [that] communicate hostile, derogatory, and/or negative slights” to any marginalized group (Sue and Spanierman, *Microaggressions in Everyday Life*, 2020)

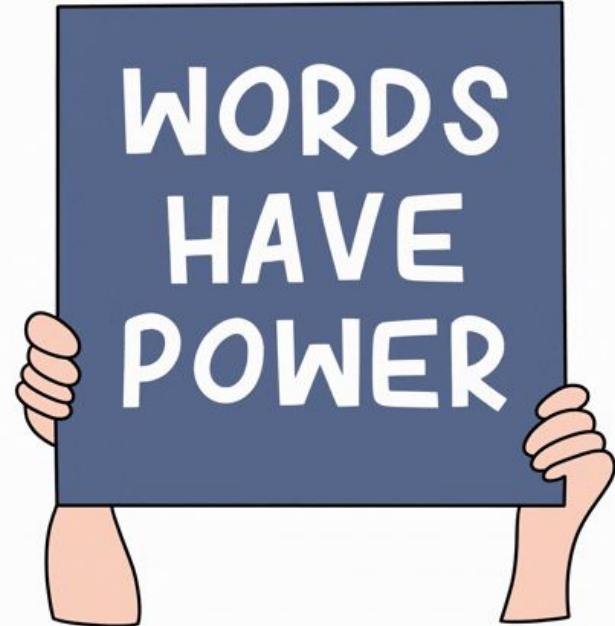


Image source:
<https://womenscenter.umbc.edu/post/114827/>

Introduction & Background

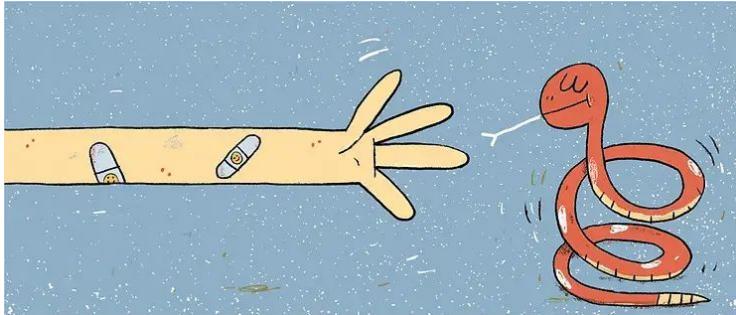


Image source:
<https://www.goethena.com/post/what-are-microaggressions/>

Examples of Microaggressions:

Asking a person of color “where are you from?”

“You’re Asian; can you help me with this math problem?”

“You’re a stay-at-home mom, so your life must be such a breeze!”

Introduction & Background

- Previous work has mostly focused on **explicit hate speech**
- Can NLP be used to classify **microaggressions**, despite their subtlety?
- **We compose BERT** and **LLM** model-based classifiers, trained on a **novel** microaggressions dataset



Image source:
<https://physicstoday.aip.org/quick-study/tip-of-the-iceberg>

Introduction & Background

- Not only **strong performance on test split**, but **strong cross-evaluation performance** on an unseen secondary dataset
- Our success criteria:
 -  Performance accuracy of **at least ten percentage points** above our baseline
 -  **High precision** on the microaggressions class

Methods: Data & Dataset Construction

- Core corpora: **selfMA** (self-reported microaggressions) & **workplaceMA** (workplace microaggressions)
- Challenge: selfMA contains **only microaggressions** → no negative class
- Early attempts: balancing with **SBIC / ToxiGen** introduced dataset artifacts
- Final solution: **LLM-generated counterfactuals** for each selfMA incident
- Resulting corpus: **selfMA_generated** – balanced micro vs non-micro, same topic & style
- Quality checks: **readability ratios** aligned with workplaceMA + **proxy audit** (style / lexical / semantic)

Methods: Models, Training & Evaluation



Image source:
Google. (2025). Gemini 2.5 Flash Image (Nano Banana version) [AI image generation model]. <https://gemini.google.com/>

- Baseline: **CNN** trained on **selfMA_generated** (binary micro vs non-micro)
- Transformers: **BERT-base**, **RoBERTa**, **HateBERT**, **DeBERTa-v3** fine-tuned on same splits
- Objective: **binary classification**, optimized with cross-entropy
- Key metrics: **accuracy**, **macro-F1**, and **recall on microaggressions**
- Cross-domain test: train on **selfMA_generated**, evaluate on **workplaceMA**
- LLM evaluation: **GPT-5.1** zero-shot with definition-based prompt on **selfMA_generated** sample

Results

- **CNN Baseline:** **78.46% accuracy** on **selfMA_generated**, but only **69.01% accuracy** on **workplace MA** → dataset-specific learning
- **Transformer Models:** Strong performance across the board
 - **Accuracies on selfMA_generated:**
Lowest: 87.69% (HateBERT)
Highest: 92.69% (DeBERTa)
 - **Accuracies on workplace MA (cross-domain):**
Lowest: 76.61% (RoBERTa)
Highest: 81.87% (DeBERTa)
 - **DeBERTa** best performer
 - **High precision and low recall** on microaggressions class in workplace MA
- **LLM (GPT-5.1):**
 - **90% accuracy** on 100-sample test
 - **Higher precision** (97.6%) and **recall** (82%) than transformer models
 - Results should be considered with some caution

Cross-domain performance of baseline and transformer models trained on selfMA_generated.				
Experimental Approach	Model	Evaluation accuracy / F1 on selfMA_generated		Cross evaluation accuracy / F1 on workplace MA
Baseline	CNN	78.46%	/ 78.45%	69.01% / 66.75 %
Single training corpus: selfMA_generated	BERT-base-cased	89.62%	/ 90 %	80.70% / 80 %
Single training corpus: selfMA_generated	hateBERT	87.69%	/ 88 %	80.12% / 80 %
Single training corpus: selfMA_generated	roBERTa-base	91.54%	/ 92 %	76.61% / 75%
Single training corpus: selfMA_generated	deBERTa-v3-base	92.69%	/ 93 %	81.87% / 81 %
Definition-based prompting classification	GPT-5.1	90%	/ 89.94 %	NA

Future Work

- **Key Achievements:**
 - Successfully demonstrated NLP models can detect microaggressions effectively
 - Created a novel dataset using LLM-generated counterfactual balancing that was crucial for semantic learning
 - Strong cross-domain performance on workplace MA dataset for transformer models
- **Impact:**
 - Reduces burden on marginalized individuals to prove microaggressions exist and cause harm
 - Enables automated detection in online spaces
- **Future Directions:**
 - Incorporate human feedback for multi-perspective dataset validation
 - Conduct fairness analysis across identity groups and microaggression types
 - Data from non-English and more diverse domains